The evidence path decides, not the model size

100 date-anchored questions on the versioned French tax code (fiscal-fr-bench); deterministic scoring, no human labels

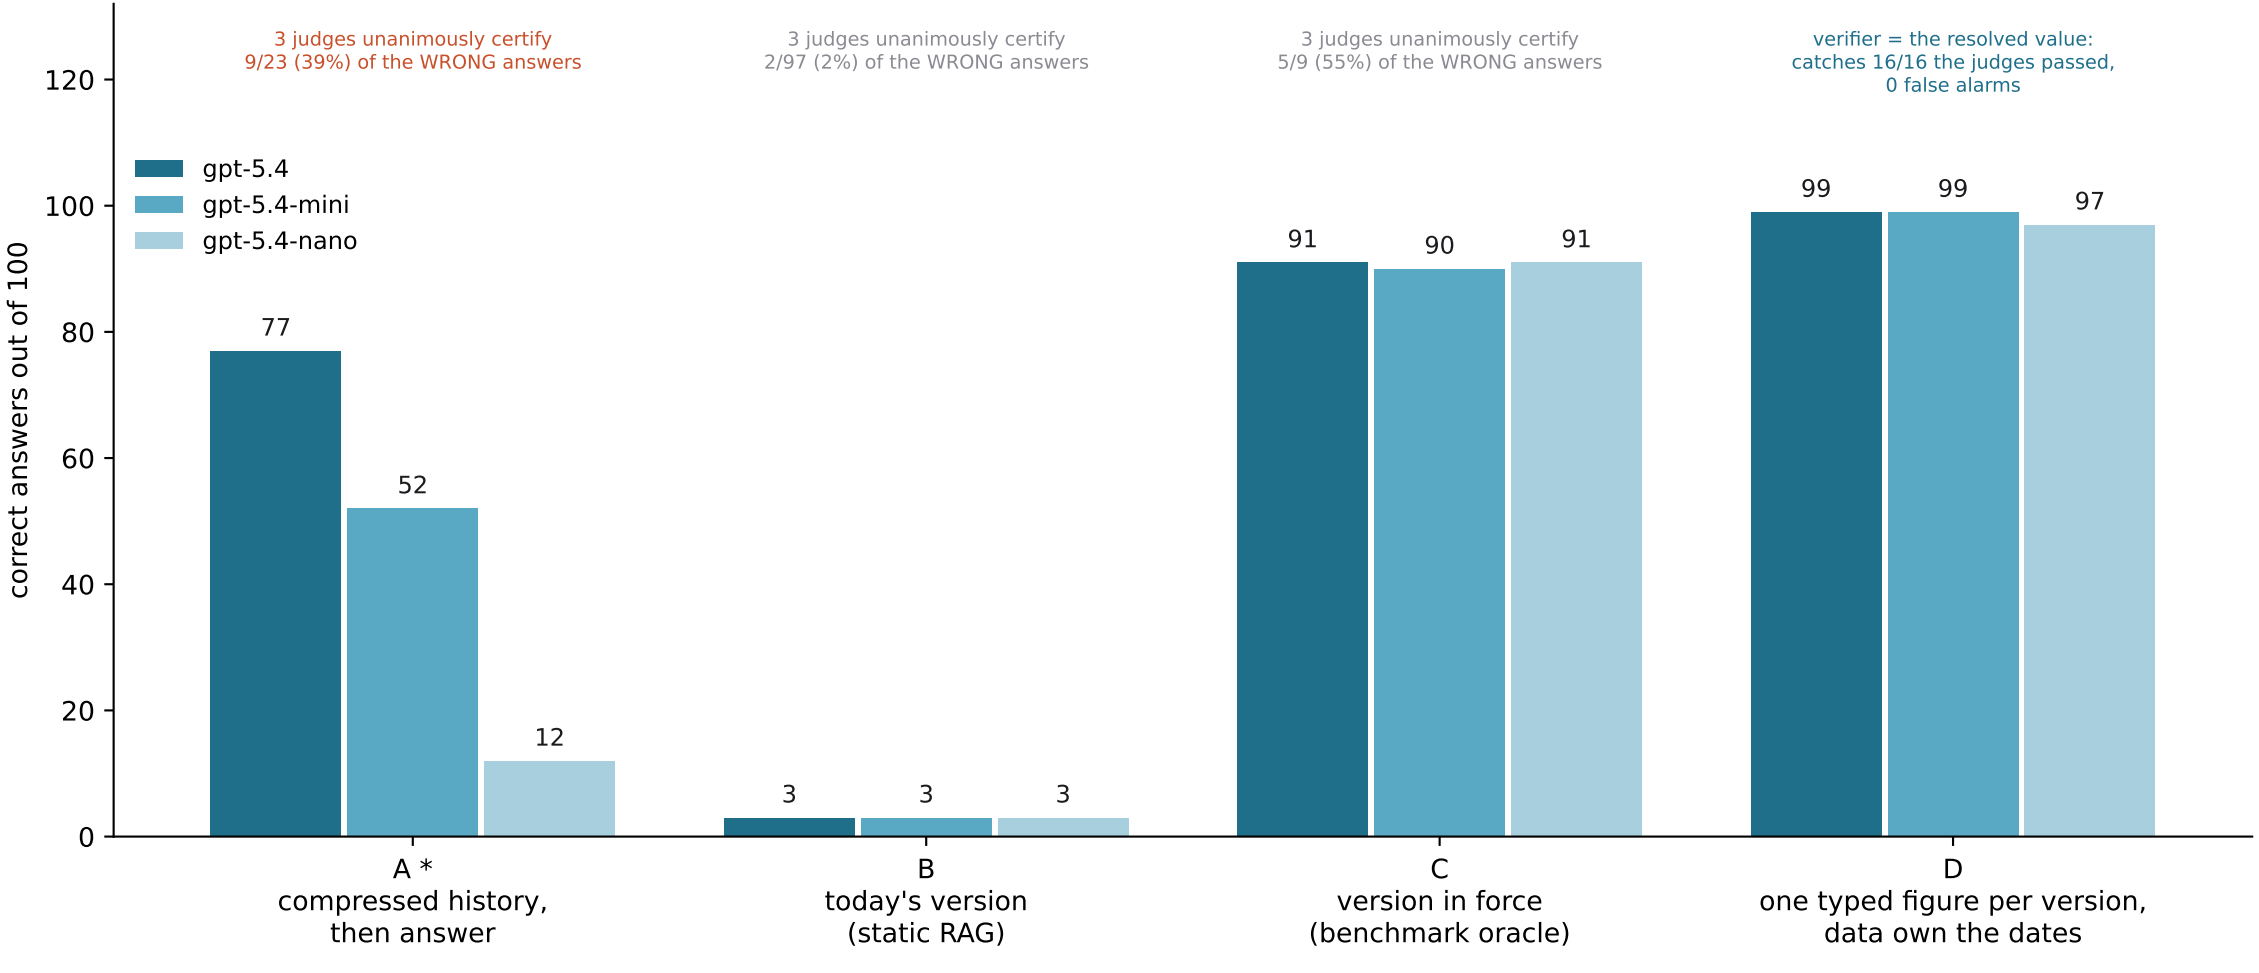

* gpt-5.4-nano: nano on A: 69 of 100 compressions failed outright; 12 of the 31 that ran were right.
Repeat runs (gpt-5.4-mini, ×3): D 99 / 99 / 99 — same single miss; C 90 / 90 / 91; B 3 / 3 / 3; A 52 / 66 / 60 — the compression step itself is stochastic.