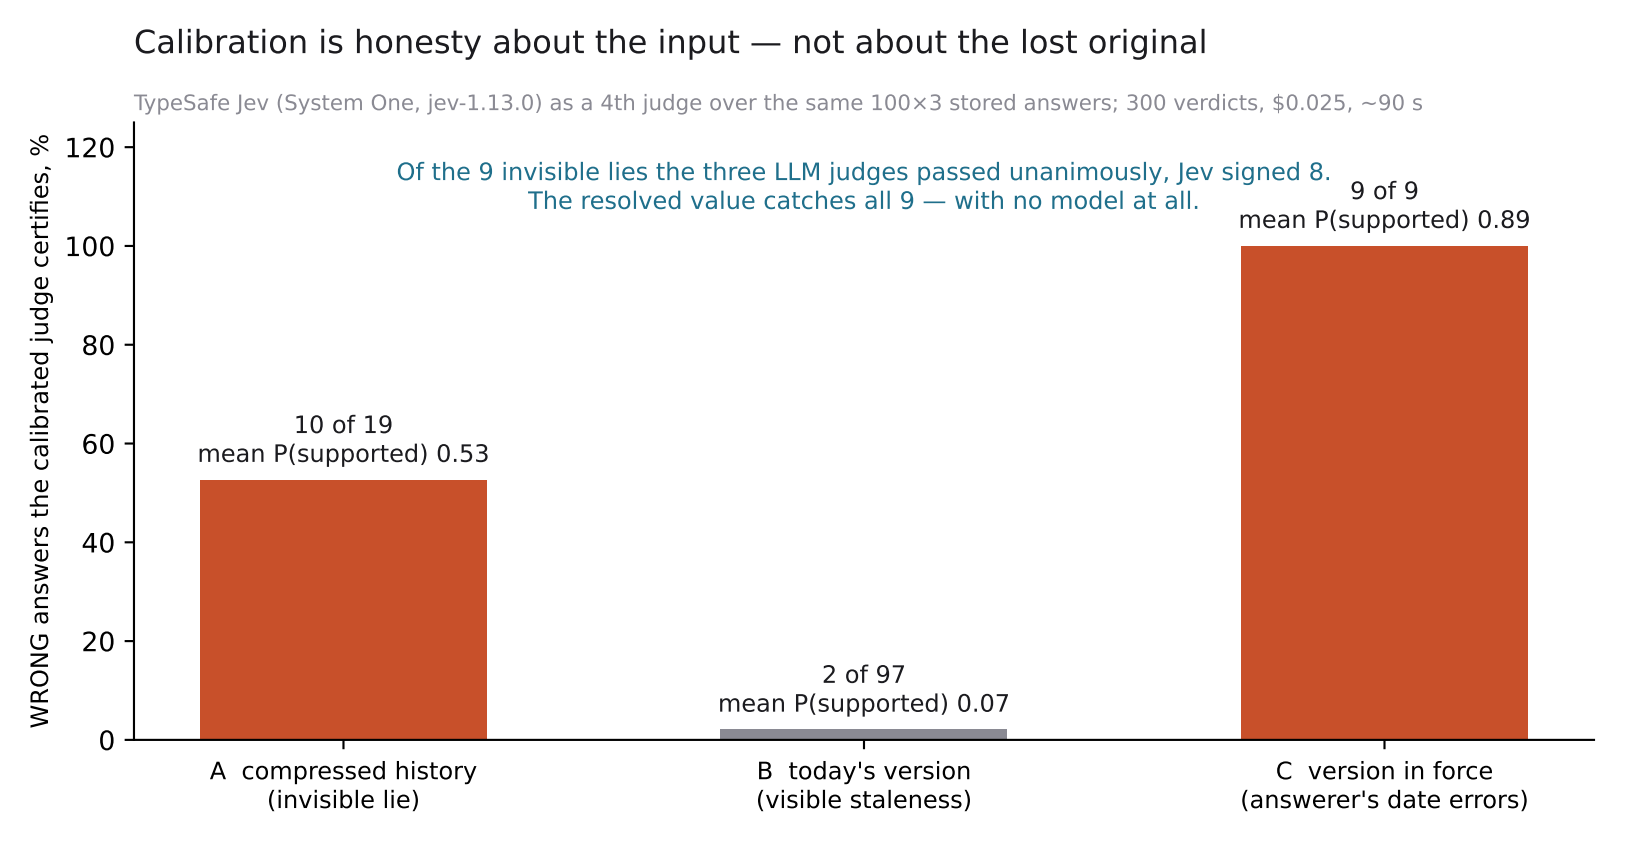
Calibration is honesty about the input — not about the lost original
TypeSafe Jev (System One, jev-1.13.0) as a 4th judge over the same 100×3 stored answers; 300 verdicts, $0.025, ~90 s
WRONG answers the calibrated judge certifies, %
120
100
80
60
40
20
0
Of the 9 invisible lies the three LLM judges passed unanimously, Jev signed 8.
The resolved value catches all 9 — with no model at all.
10 of 19
mean P(supported) 0.53
2 of 97
mean P(supported) 0.07
9 of 9
mean P(supported) 0.89
A compressed history
(invisible lie)
B today's version
(visible staleness)
C version in force
(answerer's date errors)